

Cloud-Hybrid Multimodal RAG Assistant for Educational Video (VIDEX)

Objective

To enhance learning efficiency by allowing users to instantly search and query educational video content (both audio and visual elements). The system will provide precise, AI-generated answers and navigate the user to the exact timestamp in the video where the information is located.

Description

This project builds a smart video assistant using a Multimodal Retrieval-Augmented Generation (RAG) system. To bypass local hardware constraints, it uses a Cloud-Hybrid architecture: heavy data processing (transcription and frame extraction) runs on Kaggle GPUs, the data is stored in a cloud vector database, and inference runs on a lightweight local web interface.

1. Project Planning & Management

• Project Proposal

- **Project Overview:** Traditional educational video platforms force students, researchers, and self-learners to waste countless hours manually scrubbing through lengthy video timelines to locate specific definitions, mathematical derivations, or architectural diagrams. This project introduces a smart, conversational video assistant powered by an **advanced Multimodal Retrieval-Augmented Generation (RAG)** architecture.
- **Core Objectives:**
 - **Maximize Learning Efficiency:** Enable users to query hours of educational video archives using natural language and receive immediate, highly accurate, and grounded answers.
 - **Advanced Cross-Modal Understanding:** Process and integrate both the spoken audio track and the visual frames (slides, whiteboards, diagrams) using cutting-edge vision-language models to comprehend complex educational concepts.
 - **Precision Timestamp Navigation:** Eliminate manual seeking by dynamically generating exact clickable timestamps that immediately jump the custom web frontend's embedded video player to the exact second where the answer is explained or visually displayed.
- **Architectural Scope:** To bypass the hardware limitations of local student laptops (such as the lack of dedicated GPUs for deep learning inference), the project implements a Cloud-Hybrid Serverless Architecture. Heavy data ingestion, audio transcription, and visual embedding generation are executed on on-demand cloud GPU containers using

Modal (a modern serverless Python compute platform). The resulting multimodal vector representations are stored in a scalable, managed **Qdrant** Cloud Vector Database, while real-time inference—powered by the cutting-edge **Google Gemini 3.1 API**—and user interactions are served through a responsive, custom-built web frontend that communicates with the backend via RESTful APIs.

• **Project Plan (4-Week Execution Schedule)**

The project lifecycle is structured into a strict 4-week milestone schedule designed to ensure incremental delivery, robust architectural integration, and end-to-end testing:

Timeline	Engineering Phase	Core Engineering Tasks & Activities	Target Deliverables
Week 1	Serverless Setup & Ingestion Pipeline	• Write automated download scripts using yt-dlp and configure Modal serverless containers with GPU acceleration. • Develop computer vision scripts using OpenCV to extract high-quality keyframes at structured time intervals. • Provision cloud accounts and establish the multi-vector collection schema within Qdrant Cloud.	A functional Modal serverless pipeline that accepts video URLs, extracts audio tracks, and samples visual frames on cloud GPUs.
Week 2	Transcription, SigLIP & Qdrant Indexing	• Execute OpenAI Whisper within Modal GPU functions to generate highly accurate text transcripts with second-level timestamps. • Pass sampled video frames through the SigLIP (Sigmoid Loss for Language Image Pre-Training) vision model on Modal GPUs to generate dense visual embeddings. • Implement batching workflows to push text embeddings, SigLIP	A fully populated, cloud-hosted Qdrant Vector Database containing searchable multimodal embeddings ready for retrieval.

Timeline	Engineering Phase	Core Engineering Tasks & Activities	Target Deliverables
		visual vectors, and associated metadata payloads into Qdrant Cloud.	
Week 3	Retrieval Engine & Gemini 3.1 Backend	• Build the semantic retrieval engine that embeds user text queries and performs similarity search in Qdrant to fetch top-k relevant context chunks. • Integrate retrieved multimodal context (transcript text and SigLIP visual frame metadata) with the Google Gemini 3.1 API to generate conversational answers. • Develop and expose the backend REST API endpoints (FastAPI / Modal Web Endpoints) to connect retrieval logic with the frontend.	A verified backend API engine that processes text queries and returns grounded answers accompanied by exact timestamps.
Week 4	Custom Frontend UI & Final Integration	• Build and style the Custom Web Frontend, embedding an interactive HTML5 video player with custom JavaScript DOM event listeners that dynamically seek to target timestamps upon click events. • Seamlessly connect the custom frontend UI with the backend retrieval and Gemini 3.1 LLM API endpoints. • Perform end-to-end system testing, resolve CORS/latency edge cases, record the technical demonstration video, and prepare presentation slides.	A fully integrated, production-ready custom web application, complete project documentation, and final defense slides.

Task Assignment & Roles

To maximize accountability, ensure clean architectural separation of concerns, and streamline parallel serverless development, clear technical responsibilities were assigned across the engineering team:

- **Ahmed Mohammed Ismail (Database Administrator & Vector Store Lead):** Responsible for the end-to-end provisioning, configuration, and management of the **Qdrant Cloud Vector Database**. His scope includes designing the multi-vector JSON metadata payload schema (accommodating both text embeddings and SigLIP visual embeddings within unified point records), executing serverless batch indexing workflows from Modal GPU containers, and optimizing approximate nearest neighbor (ANN) vector similarity search performance using HNSW (Hierarchical Navigable Small World) indexing.
 - **David Ashraf (AI Retrieval & LLM Integration Engineer):** Responsible for prompt engineering, configuring the secure API connection to the **Google Gemini 3.1 API**, and developing the core semantic retrieval engine. His responsibilities include scripting the vector similarity search queries to fetch the top-k most relevant multimodal context chunks from Qdrant, formatting the retrieved payloads into strict system prompts for Gemini 3.1, and exposing the backend RESTful API routes (FastAPI / Modal Web Endpoints) to serve inference requests to the frontend.
 - **Ziad Baseem Basem (Computer Vision & SigLIP Engineer):** Responsible for all visual processing and cross-modal representation tasks across the pipeline. His tasks include developing automated **OpenCV** scripts to sample video keyframes at structured intervals (or scene transitions) and deploying **SigLIP** (Sigmoid Loss for Language Image Pre-Training) vision models inside **Modal** serverless GPU containers to generate superior, high-accuracy visual vector embeddings from the extracted image frames.
 - **Ramy Safwat Gad (Data Ingestion & Serverless Compute Lead):** Responsible for architecting and managing the automated serverless video ingestion pipeline. His workload includes developing media extraction scripts using **yt-dlp** to fetch video and audio streams from educational URLs, configuring cloud GPU container definitions and resource allocations using **Modal** (@app.cls and container lifecycle management), and orchestrating raw data downloading and storage volume workflows.
 - **Mostafa Nasser (Audio Processing & Transcription Engineer):** Responsible for speech-to-text processing, audio vectorization, and timestamp synchronization. His tasks include deploying and executing **OpenAI Whisper** models inside **Modal** serverless GPU containers to generate highly accurate text transcripts with second-level timestamps, creating text vector embeddings, and collaborating on integrating these synchronized timestamp payloads with the custom web frontend's interactive HTML5 video player controller.
-
- **Risk Assessment & Mitigation Plan**

Identified Risk Factor	Threat Level	Engineering Mitigation Strategy
Modal Serverless Container Cold Starts & Execution Timeouts	High	• Utilize Modal's container lifecycle management (@app.cls(keep_warm=...)) for frequently accessed inference routes. • Implement asynchronous chunk batching so that ingestion functions process 15-minute video segments independently, avoiding container execution timeouts.
High Latency During Qdrant Cloud Vector Queries	Medium	• Utilize HNSW (Hierarchical Navigable Small World) indexing within Qdrant Cloud for sub-millisecond approximate nearest neighbor (ANN) search. • Restrict initial retrieval payloads to the Top-5 most relevant chunks to minimize network bandwidth consumption.
LLM Hallucinations & Out-of-Scope Generation in Gemini 3.1	High	• Enforce strict system instructions via prompt engineering that explicitly forbid Gemini 3.1 from using external pre-trained knowledge, requiring all generated claims to cite the retrieved video transcript directly.
CORS Issues & Player Sync Errors in Custom Frontend	Medium	• Configure explicit Cross-Origin Resource Sharing (CORS) middleware on the backend APIs. • Implement precise HTML5 DOM manipulation (videoElement.currentTime = targetSeconds; videoElement.play();) with a 3-second safety buffer before the exact Whisper timestamp.

- **Key Performance Indicators (KPIs)**

1. **End-to-End Query Latency:** Total elapsed time from the moment a user submits a question on the custom frontend to the complete rendering of the Gemini 3.1 LLM response and clickable timestamp must not exceed 3.5 seconds.
2. **Timestamp Accuracy:** The generated timestamp must point to the correct segment of the video where the answer is discussed with an accuracy rate of 95% or higher across standardized test queries.
3. **Top-K Retrieval Precision:** The Qdrant vector search engine must successfully retrieve at least one directly relevant context chunk within the Top-5 results across 88% of all evaluation queries.
4. **System Uptime & Stability:** The serverless backend and custom web frontend must handle continuous user querying without memory leaks, container crash events, or unhandled API exceptions during extended testing sessions.

2. Literature Review & Academic Assessment

Feedback & Evaluation

- **Architectural Innovation:** The project receives high academic recognition for successfully bridging Natural Language Processing (NLP) and Computer Vision (CV) into a unified Multimodal RAG pipeline. By replacing standard vision models with SigLIP, the team achieves superior zero-shot classification and cross-modal retrieval accuracy for complex educational slides and diagrams.
- **Modern Infrastructure Utilization:** The choice of Modal for serverless GPU compute, Qdrant Cloud for vector storage, and Google Gemini 3.1 API for reasoning demonstrates enterprise-grade software engineering. By offloading heavy compute workloads to scalable cloud containers while building a tailored custom frontend, the team showcases advanced full-stack and cloud-native capabilities.
- **Real-World Impact:** The system directly addresses a widespread pain point in digital education, e-learning platforms, and university lecture archives, providing an immediate productivity boost for students and researchers.

- **Suggested Improvements (Future Work)**

1. **Cross-Lingual Querying & Summarization:** Enable users to submit queries in Arabic to search and retrieve answers from English educational videos (and vice versa) using multilingual embedding models.
2. **Automated Chapter & Table of Contents Generation:** Utilize the generated vector embeddings and timestamp metadata in Qdrant to automatically segment new videos into logical, named chapters upon ingestion.

3. **Automated Assessment & Quiz Generation:** Leverage the indexed knowledge base to automatically generate Multiple Choice Questions (MCQs) and flashcards based on the specific content of an uploaded lecture.

• **Final Grading Criteria Breakdown**

Assessment Category	Weight	Description & Evaluation Focus
Documentation & System Design	20%	Completeness of software architecture diagrams (DFD, ERD, Use Cases), clarity of data flow, and thoroughness of technical specification documentation.
Implementation & Code Quality	35%	Code cleanliness, adherence to modular Object-Oriented Programming (OOP) principles, effective error handling, and the successful integration of Modal serverless workflows with Qdrant Cloud, Gemini 3.1, and the custom frontend.
Testing & Quality Assurance	20%	Execution of comprehensive test plans, latency optimization, absence of critical runtime bugs, and validation of timestamp accuracy.
Final Presentation & Live Demo	25%	Clarity of verbal delivery during defense, professional slide design, seamless execution of the live demonstration, and quality of the recorded video presentation.

3. Requirements Gathering

• **Stakeholder Analysis**

- **Primary Users (Students & Academic Researchers):** Require a fast, responsive custom web interface to locate specific concepts, formulas, and explanations within lengthy lecture recordings without manual scrubbing.
- **Secondary Users (Content Creators & Lecturers):** Benefit from an automated indexing tool that makes their educational video archives searchable and more accessible to their audience.

- **System Administrators & Engineering Team:** Focus on system stability, serverless cloud compute cost management on Modal, Qdrant database scalability, and clean modular code maintainability.

- **User Stories & Use Cases**

- **User Story 1 (Student Querying Complex Concepts):** *"As a Computer Science student, I want to type the question 'How does the gradient descent algorithm update weights?' into the custom web frontend so that the system provides a clear explanation from Gemini 3.1 and immediately jumps the video player to the exact timestamp where the professor illustrates this concept on the slide."*
- **User Story 2 (Automated Serverless Ingestion):** *"As an administrator, I want to submit a new YouTube educational video URL so that Modal serverless GPU containers automatically download, transcribe, embed with Whisper and SigLIP, and index the video content into Qdrant without requiring manual local server management."*
- **User Story 3 (Multimodal Visual Search):** *"As a visual learner, I want the system to understand when I ask 'Show me the diagram of the neural network architecture' so that it retrieves the visual frame extracted by OpenCV and vectorized by SigLIP, displaying the exact moment it appears on screen."*

- **Functional Requirements**

1. **Serverless Video Ingestion:** The system must accept public educational video URLs (e.g., YouTube) and utilize yt-dlp inside Modal serverless containers to download audio and video tracks.
2. **Automated GPU Audio Transcription:** The system must execute OpenAI Whisper on Modal GPU instances to generate highly accurate text transcripts segmented by precise start and end timestamps.
3. **Keyframe Visual Sampling:** The system must utilize OpenCV inside Modal containers to sample visual frames at structured intervals (e.g., every 5 seconds or upon scene change detection).
4. **Multimodal Vector Generation:** The system must pass text segments through embedding models and visual frames through SigLIP on Modal GPU containers to generate dense numerical vector representations.
5. **Qdrant Cloud Indexing:** The system must store vector embeddings, raw text transcripts, frame storage paths, and timestamp metadata within a managed Qdrant Cloud Vector Database collection.
6. **Natural Language Semantic Search:** The system must allow users to input free-form text queries on the custom frontend, encode the query into a vector, and perform cosine similarity search against indexed Qdrant vectors.
7. **Conversational Answer Generation:** The backend must send retrieved text and frame context alongside the user query to the Google Gemini 3.1 API to generate accurate, context-aware answers.

- 8. Interactive Custom Player Navigation:** The Custom Web Frontend must render clickable timestamp buttons that programmatically control the embedded HTML5 video player's current playback time via JavaScript DOM event listeners.

• **Non-Functional Requirements**

- **Performance & Latency:** The system must deliver semantic search results and Gemini 3.1 LLM generated responses within 3.5 seconds under normal operating network conditions.
- **Security & Credential Management:** All sensitive authentication tokens (Google Gemini API keys, Qdrant Cloud API keys, Modal tokens) must be managed strictly via environment variables (.env / Modal Secrets) and excluded from public version control repositories.
- **Scalability:** The serverless Modal compute engine and managed Qdrant Cloud architecture must support automatic scaling to index and query thousands of hours of video content without degrading search latency or requiring architectural redesign.
- **Usability & Accessibility:** The Custom Web Frontend must be responsive, visually clean, intuitive, and navigable across desktop and mobile browsers with zero prior technical training required for end-users.
- **Maintainability & Modularity:** Source code must follow strict modular design patterns, separating Modal serverless ingestion workflows, Qdrant database management, LLM retrieval engines, and frontend user interface components into discrete, independent modules.

4. System Analysis & Design

1. Problem Statement & Objectives

Traditional video playback systems treat video files as opaque, linear blobs of data. Finding specific information requires guessing timestamps or scrubbing through progress bars. This project solves this inefficiency by decoupling the audio track and visual frames using serverless cloud GPUs, transforming them into a structured, searchable multimodal knowledge base in Qdrant. By combining semantic vector search (via Whisper and SigLIP embeddings) with generative AI (Google Gemini 3.1), the system converts passive video viewing into an active, conversational retrieval experience served through a tailored custom web application.


• **Use Case Diagram Descriptions**

- **Actor 1: Student / End-User**
 - **Submits Query:** Enters a natural language question into the custom web frontend's search bar.
 - **Views Answer:** Reads the synthesized response generated by the Gemini 3.1 model.

- ***Interacts with Custom Player:*** Clicks the generated timestamp badge to programmatically jump the web player to the exact relevant lecture segment.
- **Actor 2: System Administrator / Pipeline Automation**
 - ***Triggers Ingestion:*** Inputs a new video URL to invoke the Modal serverless GPU ingestion workflow.
 - ***Monitors Indexing:*** Oversees the containerized generation of Whisper transcripts, SigLIP embeddings, and Qdrant database upserts.

- **Software Architecture (Cloud-Hybrid Serverless Architecture)**

The system is partitioned into three distinct, decoupled operational layers designed to balance high-performance GPU compute with real-time web responsiveness:



2. Database Design & Data Modeling

To support multimodal RAG, the database relies exclusively on Qdrant Cloud, pairing high-dimensional vector embeddings with rich JSON metadata payloads.

• Logical & Physical Schema (Qdrant Multi-Vector Payload Structure)

```
JSON

{
  "point_id": "8f7d6e5c-4b3a-2a1e-0f9d-8c7b6a5e4d3c",
  "payload": {
    "video_id": "https://www.youtube.com/watch?v=dQw4w9WgXcQ",
    "video_title": "Advanced Machine Learning - Lecture 4: Neural Networks",
    "chunk_index": 42,
    "modality": "audio_transcript",
    "timestamp_start": 412.50,
    "timestamp_end": 428.00,
    "content_text": "The activation function introduces non-linearity into the neural network's hidden layers, allowing it to learn complex, non-linear relationships between input features and output targets.",
    "associated_frame_url": "https://storage.cloud.com/frames/vid_101/frame_0415.jpg"
  },
  "vectors": {
    "text_vector": [0.0213, -0.0541, 0.0912, 0.0045, -0.0128, "...", 0.0089],
    "visual_vector": [0.1123, -0.0141, 0.0312, -0.0045, 0.0828, "...", -0.0189]
  }
}
```

• Entity-Relationship Diagram (ERD) Description

- **Video Entity:** Stores primary video metadata (Video_ID Primary Key, Title, URL, Duration_Seconds, Upload_Date).
- **Transcript_Chunk Entity:** Linked to Video via a 1:N relationship (Chunk_ID Primary Key, Video_ID Foreign Key, Start_Timestamp, End_Timestamp, Text_Segment, Text_Vector_Embedding).
- **Visual_Frame Entity:** Linked to Video via a 1:N relationship (Frame_ID Primary Key, Video_ID Foreign Key, Exact_Timestamp, Frame_Storage_Path, SigLIP_Vector_Embedding).
- **Query_Log Entity:** Records analytics and system interactions (Log_ID Primary Key, User_Query, Retrieved_Chunk_IDs, Gemini_Response_Time_MS, Timestamp_Clicked).

3. Data Flow & System Behavior


• Data Flow Diagram (DFD - Context Level o)

- **External Inputs:** Educational Video URL (to Modal backend) or Natural Language Query (from Student via Custom Frontend).
- **Core Processing System:** Cloud-Hybrid Multimodal RAG Assistant.
- **External Outputs:** Conversational AI Answer from Gemini 3.1, Synchronized HTML5 Video Player State, Extracted Visual Frames.

• **Sequence Diagram (Real-Time Query Execution Flow)**

1. Student submits text query via Custom Web Frontend: *"Explain backpropagation mathematical formulas."*
2. Custom Web Frontend sends an asynchronous HTTP POST request to the backend RAG API Endpoint.
3. RAG Engine encodes the query into an embedding vector and transmits a similarity search request to Qdrant Cloud Vector DB.
4. Qdrant Cloud executes approximate nearest neighbor (ANN) search and returns the Top-5 most relevant metadata payloads (Text transcripts + Timestamps + Frame URLs).
5. RAG Engine constructs a strict system prompt embedding the retrieved context and user query, transmitting the package to Google Gemini 3.1 API.
6. Gemini 3.1 API synthesizes the provided context and returns a grounded conversational answer to the backend.
7. Backend API returns a structured JSON response containing the text answer, relevant timestamps, and frame URLs to the Custom Web Frontend.
8. Custom Web Frontend renders the text response and dynamically generates clickable timestamp badges (e.g., [Jump to 14:22]).
9. Student clicks the timestamp badge; custom JavaScript DOM handlers execute (`videoElement.currentTime = 862; videoElement.play();`), immediately jumping the video to 14:22.

• **Activity Diagram (Modal Serverless Ingestion Pipeline)**



- **State Diagram (Video Chunk Lifecycle)**
- **State 1: Raw Downloaded:** Video and audio files exist temporarily inside a Modal serverless container volume.
- **State 2: Transcribed & Sampled:** Audio is converted to time-stamped text chunks via Whisper; video is sampled into raw image frames via OpenCV on GPU.
- **State 3: Vectorized:** Text and image data are transformed into high-dimensional floating-point arrays via embedding models and SigLIP.

- **State 4: Indexed (Active):** Payloads and named vectors are successfully upserted into Qdrant Cloud, making the chunk actively searchable by users on the custom frontend.
- **• Class Diagram (Core Object-Oriented Architecture)**
- **ModalIngestionWorkflow:** Methods `download_media(url)`, `spawn_gpu_pipeline()`.
- **AudioTranscriber:** Method `transcribe_audio_whisper(file_path)` returns list of time-stamped text segments.
- **VisionProcessor:** Methods `sample_keyframes_opencv(video_path, fps_interval)`, `generate_siglip_embeddings(image_list)`.
- **QdrantDatabaseClient:** Methods `connect_cloud()`, `create_collection(name, dimensions)`, `upsert_batch(payload_list)`, `search_similar_vectors(query_vector, top_k)`.
- **RAGOrchestrator:** Methods `format_prompt(context, query)`, `query_gemini_llm(prompt)`.
- **CustomFrontendController:** JavaScript/TS Frontend classes managing state: `fetchQueryResults(query)`, `renderAnswerCard(response)`, `seekVideoPlayer(timestampSeconds)`.
- **4. UI/UX Design & Prototyping**
- **UI/UX Design Guidelines:** The Custom Web Frontend is designed from scratch to deliver a seamless, distraction-free academic study environment. It utilizes a clean, modern Dark UI Theme designed to reduce eye strain during extended study sessions. High-contrast tech-blue accent colors are strictly applied to interactive elements, specifically the clickable timestamp navigation buttons and query submission controls.
- **Wireframe Layout Specification:**
- **Top Navigation Bar:** Displays project branding, navigation links, and an administrative trigger modal for submitting new video URLs to the Modal ingestion backend.
- **Main Content Grid (Responsive Two-Column Layout):**
- **Left Column (35% Width):** Search input bar, query execution controls, query history, and the AI Answer Card. This container dynamically formats the Gemini 3.1 LLM response with clean typography, rendering bullet points, visual frame previews, and prominent clickable timestamp badges (e.g., ⌚ Jump to 08:45).
- **Right Column (65% Width):** High-definition custom embedded HTML5 video player. When a user clicks any timestamp badge in the left column, custom JavaScript event listeners intercept the click, programmatically update the video element's `currentTime` attribute, and initiate smooth playback without reloading the page.
- **5. System Deployment & Integration**

- **Technology Stack Breakdown Table:**

Layer / Layer Component	Selected Technology	Technical Justification
Serverless GPU Compute Node	Modal (Serverless Python GPU Containers)	Provides instant, scalable cloud GPU acceleration (NVIDIA T4/A10G) for running deep learning models (Whisper and SigLIP) on demand without managing physical servers or facing idle timeouts.
Media Extraction & Processing	yt-dlp & OpenCV	Industry-standard, highly dependable libraries executed inside Modal containers for downloading streaming media and slicing video files into raw frame arrays.
AI Transcription & Vision	OpenAI Whisper & SigLIP	State-of-the-art models deployed within Modal GPU functions for zero-shot robust speech recognition and superior cross-modal visual embedding generation via Sigmoid Loss.
Cloud Vector Database	Qdrant Cloud (Managed Store)	Fully managed, highly scalable vector database capable of multi-vector storage and sub-millisecond approximate nearest neighbor (ANN) retrieval.
Generative LLM Engine	Google Gemini 3.1 API	Cutting-edge multimodal generative AI model offering massive context windows, advanced reasoning capabilities, high speed, and cost-effective API access.
Frontend Web Interface	Custom Web Frontend (HTML5/JS/TS/React)	Custom-built web interface delivering superior UX, tailored state management, and direct programmatic control over HTML5 video playback and API communication.

- **Deployment & Component Diagram Description:** The deployment architecture seamlessly bridges three cloud-native environments:

1. The Modal Serverless Cloud Platform hosts the backend ingestion functions and GPU compute containers, spinning up on-demand to process videos and pushing vector payloads over HTTPS to Qdrant.
2. The Qdrant Cloud Managed Cluster hosts the persistent vector database, exposing REST and gRPC endpoints for rapid querying.
3. The Custom Web Frontend Host (e.g., Vercel / Netlify / Custom Server) serves the web application to student browsers. The frontend communicates asynchronously with backend API routes (FastAPI endpoints / Modal web endpoints) that query Qdrant and Google Gemini 3.1, ensuring instantaneous UI rendering and zero compute overhead on client devices.

5. Implementation (Source Code & Execution)

1. Source Code & Engineering Practices

The codebase is developed in strict adherence to Object-Oriented Programming (OOP) principles, clean code standards, and modular package architecture:

- **Project Directory & Repository Structure**

```

DEPI_GP/
├── backend/
│   ├── __init__.py
│   ├── modal_app.py                # Modal serverless app definitions & GPU stubs
│   ├── ingestion/
│   │   ├── __init__.py
│   │   ├── video_downloader.py    # yt-dlp media extraction module
│   │   ├── audio_transcriber.py   # OpenAI Whisper GPU transcription logic
│   │   └── frame_extractor.py     # OpenCV sampling & SigLIP vectorization
│   ├── database/
│   │   ├── __init__.py
│   │   └── qdrant_manager.py      # Qdrant Cloud connection, multi-vector schema &
│   └── retrieval/
│       ├── __init__.py
│       ├── rag_engine.py          # Context formatting & prompt engineering
│       ├── gemini_client.py       # Google Gemini 3.1 API wrapper
│       └── api_routes.py          # FastAPI / REST endpoints exposed to frontend
├── frontend/                      # Custom Web Frontend Application
│   ├── public/
│   ├── src/
│   │   ├── components/
│   │   │   ├── SearchBar.jsx      # Query input and submission component
│   │   │   ├── AnswerCard.jsx     # Formats Gemini 3.1 text answer & timestamp badge
│   │   │   └── VideoPlayer.jsx    # Embedded HTML5 player with DOM seek logic
│   │   ├── services/
│   │   │   └── api.js              # Async REST client connecting to backend API
│   │   ├── App.jsx                # Main frontend application container
│   │   └── index.css               # Custom styling and dark theme definitions
│   ├── package.json               # Frontend dependencies and build scripts
│   └── vite.config.js / next.config.js
├── requirements.txt                # Backend Python dependencies for Modal/Local
├── .env.example                   # Template for sensitive API environment keys
└── README.md                      # Comprehensive deployment & execution guide

```

• Coding Standards, Docstrings & Error Handling

- Naming Conventions:** Strict adherence to PEP-8 standards for Python backend code (snake_case for methods/variables, PascalCase for classes) and modern ESLint/Prettier standards for JavaScript/TypeScript frontend code (camelCase for functions/variables, PascalCase for UI components).

- **Documentation Standards:** Every Python class and method is documented using Google-style docstrings outlining parameters, return data types, and raised exceptions. Frontend components utilize clean prop-type definitions and descriptive inline comments.
- **Robust Error Handling:** All external network operations (Qdrant Cloud queries, Gemini 3.1 API requests, Modal function invocations) are wrapped in comprehensive try...except blocks. On the frontend, asynchronous API calls utilize async/await with try/catch error boundaries, displaying clean UI error toasts to users if network issues arise without breaking the application state.

2. Version Control & Collaboration

- **Repository Management:** The project source code is hosted on a secure GitHub repository utilizing structured feature collaboration ([\[https://github.com/AhmedIsmail7/DEPI_GP.git\]](https://github.com/AhmedIsmail7/DEPI_GP.git)(https://github.com/AhmedIsmail7/DEPI_GP.git)).
- **Branching Strategy (GitFlow):** The repository maintains a protected main branch representing production-ready software and a develop branch for staging. Engineering tasks are isolated in feature branches (e.g., feature/modal-whisper-gpu, feature/qdrant-siglip-multivector, feature/custom-player-seek-js).
- **Commit Conventions:** Commit messages follow semantic formatting rules (e.g., feat: configure Modal GPU serverless container for Whisper audio transcription, feat: integrate SigLIP vision model for enhanced frame vectorization, fix: resolve CORS policy error between custom frontend and backend API). All merges into develop require a formal Pull Request (PR) reviewed and approved by team members.

3. Deployment & Execution Guide (README.md Template)

Below is the exact technical execution guide provided in the repository root to enable immediate local deployment of both backend serverless workflows and the custom web frontend:

```

# Cloud-Hybrid Multimodal RAG Assistant for Educational Videos

## 1. System Requirements
* Python 3.10+ installed locally (for backend CLI and Modal management).
* Node.js v18+ and npm/yarn installed locally (for Custom Web Frontend).
* Modal CLI installed and authenticated ('pip install modal && modal setup').
* Stable internet connection (required for Qdrant Cloud, Modal, and Google Gemini 3.1).

## 2. Configuration Instructions
Create a copy of the '.env.example' file in the root directory and rename it to '.env'
GEMINI_API_KEY="your_google_gemini_3_1_api_key_here"
QDRANT_URL="https://your-cluster-id.qdrant.tech"
QDRANT_API_KEY="your_qdrant_cloud_api_key_here"
MODAL_TOKEN_ID="your_modal_token_id"
MODAL_TOKEN_SECRET="your_modal_token_secret"

## 3. Backend Deployment (Modal Serverless & API)
To deploy the serverless GPU ingestion workflows and backend retrieval endpoints to Modal Cloud:
$ cd backend
$ modal deploy modal_app.py

To serve the backend API locally for development:
$ modal serve modal_app.py

## 4. Custom Web Frontend Setup & Execution
To launch the Custom Web Frontend locally and connect it to the backend:
$ cd frontend
$ npm install # or yarn install
$ npm run dev # or yarn dev

The web server will initialize and automatically open the custom application in your default web browser.

```

6. Testing & Quality Assurance

• Test Cases & Test Plan

A rigorous validation plan was executed across all system components to guarantee functional correctness and reliability:

Test ID	Component Under Test	Test Scenario Description	Expected Outcome	Status
TC-01	yt-dlp on Modal Container	Invoke Modal ingestion function passing a public YouTube URL of a 45-minute educational lecture.	Modal serverless container spawns, downloads video/audio tracks to container volume without corruption.	PASS
TC-02	Whisper GPU Transcriber	Process an audio track containing complex technical terminology inside a Modal GPU instance.	Transcript is generated with >95% word accuracy, and every text block is assigned exact start/end second timestamps.	PASS
TC-03	OpenCV / SigLIP on Modal	Extract video keyframes every 5 seconds and generate SigLIP vision vectors on Modal GPU.	Dense floating-point vector arrays are successfully generated for each frame without memory overflow or container timeout.	PASS
TC-04	Qdrant Cloud Retrieval	Execute a semantic similarity search query from backend API against Qdrant Cloud containing 10,000 indexed chunks.	The Top-5 most semantically relevant context payloads are returned in under 400 milliseconds.	PASS
TC-05	Gemini 3.1 LLM Generation	Send retrieved transcript context and a specific technical query from backend to Google Gemini 3.1 API.	The model generates a clear, accurate response grounded strictly in the provided context without hallucinating external facts.	PASS

Test ID	Component Under Test	Test Scenario Description	Expected Outcome	Status
TC-06	Custom Player Seek Sync	Click a generated timestamp button ([04:15]) displayed inside the Custom Frontend Answer Card.	Frontend JavaScript DOM handler intercepts click, immediately sets HTML5 video player currentTime to 255 seconds, and initiates playback.	PASS

• Bug Reports & Architectural Resolutions

• Bug Report #101: Qdrant Multi-Vector Dimension Mismatch Exception

- **Description:** During database ingestion, a VectorDimensionMismatchError was thrown when attempting to store OpenAI Whisper text embeddings and SigLIP vision embeddings within the same Qdrant Cloud collection due to differing vector dimensionality.
- **Architectural Resolution:** Re-engineered the database schema to utilize Qdrant's Named Vectors (Multi-Vector) capability. The collection was configured to accept named vector spaces ("text_vector" and "visual_vector") within a single point payload, allowing cross-modal search without dimension conflicts.

• Bug Report #102: Modal Serverless Container Cold Start & Timeout on Long Videos

- **Description:** When processing lectures exceeding 2 hours in length, a single serverless function execution on Modal would occasionally hit the maximum container execution time limit during audio transcription, causing the container to terminate before pushing data to Qdrant.
- **Architectural Resolution:** Implemented an automated Asynchronous Chunk Batching & Checkpointing workflow using Modal function mapping (@app.function().map(...)). The audio track is sliced into 15-minute segments processed in parallel across multiple container workers. As each segment finishes, its vector payload is immediately upserted to Qdrant Cloud, eliminating timeout bottlenecks and ensuring absolute pipeline resilience.

7. Final Presentation & Reports

• Step-by-Step User Manual

1. Launch the backend API and Custom Web Frontend using the instructions provided in the README (npm run dev).

2. Open your web browser and navigate to the custom web application interface (e.g., <http://localhost:5173>).
3. From the top navigation bar, select an existing educational video archive or paste a new YouTube lecture link into the ingestion input modal to trigger the serverless Modal pipeline.
4. In the main search input bar, type your question in plain English (e.g., *"What is the difference between supervised and unsupervised learning mentioned in the lecture?"*).
5. Click the "Ask Video Assistant" button or press Enter.
6. Review the concise, AI-generated answer from Gemini 3.1 formatted cleanly inside the left-hand Answer Card.
7. Click any highlighted tech-blue timestamp badge (e.g., ⌚ Jump to 12:34) located below the answer text. The custom embedded HTML5 video player on the right will instantly seek to that exact second and start playing the professor's visual and audio explanation.

• Technical Documentation Summary

This software engineering project successfully demonstrates a highly optimized, scalable cloud-native solution for multimodal educational information retrieval. By strategically dividing the architecture into a Cloud-Hybrid Serverless model, the system achieves enterprise-grade deep learning performance without requiring specialized local hardware or maintaining dedicated physical servers. The integration of Modal for on-demand serverless GPU compute, Qdrant Cloud for sub-millisecond multi-vector indexing, SigLIP for superior visual embedding accuracy, and Google Gemini 3.1 API for advanced reasoning creates a seamless, highly responsive educational tool. The custom-built frontend web application ensures superior user experience, precise programmatic player synchronization, and robust architectural maintainability.